

Convolutional Neural Networks

This notebook explains CNNs through the implementation in sklearn.

We train a CNN to classify images from the [CIFAR-10 dataset](#).

- The CIFAR-10 dataset consists of 60,000 colour images of size (32x32) in 10 classes, with 6000 images per class.
- Here are the classes in the dataset, as well as 10 random images from each:
 - airplane
 - automobile
 - bird
 - cat
 - deer
 - dog
 - frog
 - horse
 - ship
 - truck

1. Load CIFAR-10 Database

```
In [27]: import keras
from keras.datasets import cifar10

# load the pre-shuffled train and test data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

2. Visualize Some Training Images

```
In [28]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib inline

fig = plt.figure(figsize=(20,5))
for i in range(40):
    ax = fig.add_subplot(4, 10, i + 1, ticksa=[], yticks=[])
    ax.imshow(np.squeeze(x_train[i]))
```



3. Rescale the Images by Dividing Every Pixel in Every Image by 255

- Let's see how a datapoint looks like:

```
In [29]: print(x_train[1])

[[[154 177 187]
 [126 137 136]
 [105 104 355]
 ...
 [ 91 95 71]
 [ 87 90 71]
 [ 79 81 70]]

 [[140 160 169]
 [145 153 154]
 [125 125 118]
 ...
 [ 96 99 78]
 [ 77 80 62]
 [ 71 73 61]]

 [[140 155 161]
 [139 146 149]
 [115 115 112]
 ...
 [ 79 82 64]
 [ 68 70 55]
 [ 67 69 55]]

 ...

 [[175 167 166]
 [156 154 160]
 [154 160 170]
 ...
 [ 42 34 36]
 [ 61 53 57]
 [ 93 83 91]]

 [[165 154 128]
 [156 152 130]
 [159 161 142]
 ...
 [103 93 96]
 [123 114 120]
 [131 121 131]]

 [[163 148 120]
 [158 148 122]
 [163 156 133]
 ...
 [143 133 139]
 [143 134 142]
 [143 133 144]]]

...

 [[175 167 166]
 [156 154 160]
 [154 160 170]
 ...
 [ 42 34 36]
 [ 61 53 57]
 [ 93 83 91]]

 [[165 154 128]
 [156 152 130]
 [159 161 142]
 ...
 [103 93 96]
 [123 114 120]
 [131 121 131]]

 [[163 148 120]
 [158 148 122]
 [163 156 133]
 ...
 [143 133 139]
 [143 134 142]
 [143 133 144]]]
```

- Let's normalize each cell value [0,255], so that each cell value lies within the same distribution [0,1].
- That is, the values change from [0,255] → [0,1].
- We normalize using the formula:

$$\frac{X - X_{min}}{X_{max} - X_{min}}$$

- Here, $X_{min}=0$ and $X_{max}=255$.
- Hence,

$$\frac{X - X_{min}}{X_{max} - X_{min}} = \frac{X}{255}$$

```
In [30]: x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255
```

4. Break Dataset into Training, Testing, and Validation Sets

- A datapoint in the dataset could belong to any one of the 10 classes mentioned earlier.
- Let's convert the output into one-hot encoded vector of size 10.
- This will be helpful for our CNN to recognize a given image datapoint into one of the 10 classes.

```
In [31]: from keras.utils import np_utils

# one-hot encoding
num_classes = len(np.unique(y_train))
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [32]: # break training set into training and validation sets
(x_train, y_valid) = x_train[5000:], x_train[:5000]
(y_train, y_valid) = y_train[5000:], y_train[:5000]

# print shape of training set
print("x_train shape:", x_train.shape)

# print number of training, validation, and test images
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
print(x_valid.shape[0], 'validation samples')

x_train shape: (45000, 32, 32, 3)
45000 train samples
10000 test samples
5000 validation samples
```

5. Data Augmentation

- Data augmentation helps in creating invariance in the data, especially in image data.
- Invariance in data could be for:
 - Shift
 - Rotation
 - Zoom
 - Noise, etc

```
In [33]: from keras.preprocessing.image import ImageDataGenerator

# create and configure augmented image generator
datagen_train = ImageDataGenerator(
    width_shift_range=0.1, # randomly shift images horizontally (10% of total width)
    height_shift_range=0.1, # randomly shift images vertically (10% of total height)
    horizontal_flip=True) # randomly flip images horizontally

# fit augmented image generator on data
datagen_train.fit(x_train)
```

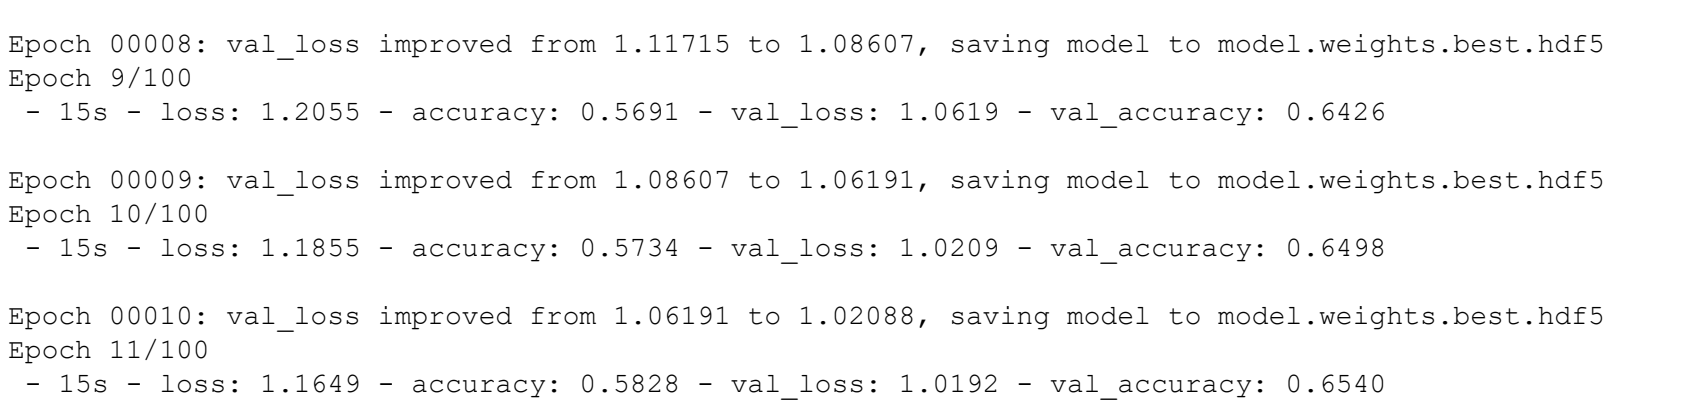
6. Visualize Original and Augmented Images

```
In [34]: import matplotlib.pyplot as plt

# take subset of training data
x_train_subset = x_train[:12]

# visualize subset of training data
fig = plt.figure(figsize=(20,2))
for i in range(0, len(x_train_subset)):
    ax = fig.add_subplot(1, 12, i+1)
    ax.imshow(x_train_subset[i])
fig.savefig('Subset of Original Training Images', fontsize=20)
plt.show()

# visualize augmented images
fig = plt.figure(figsize=(20,2))
for i in range(0, len(x_train_subset)):
    ax = fig.add_subplot(1, 12, i+1)
    ax.imshow(x_train_subset[i])
fig.savefig('Augmented Images', fontsize=20)
plt.show()
break;
```



7. Define the Model Architecture

```
In [35]: # Necessary model parameters:
input_dim = 10
input_shape = x_train[1].shape
batch_size = 30
epochs = 100
```

```
In [36]: from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

model = Sequential()
```

```
model.add(Conv2D(filters=64, kernel_size=3, padding='same', kernel_initializer='glorot_uniform',
activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2, padding='same', kernel_initializer='glorot_uniform',
activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.5))
model.add(Conv2D(filters=32, kernel_size=2, kernel_initializer='glorot_uniform', padding='same',
activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(512, kernel_initializer='glorot_uniform', activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 32, 32, 64)	1792
max_pooling2d_10 (MaxPooling)	(None, 16, 16, 64)	0
conv2d_11 (Conv2D)	(None, 16, 16, 32)	8224
max_pooling2d_11 (MaxPooling)	(None, 8, 8, 32)	0
dropout_10 (Dropout)	(None, 8, 8, 32)	0
conv2d_12 (Conv2D)	(None, 8, 8, 32)	4128
max_pooling2d_12 (MaxPooling)	(None, 4, 4, 32)	0
dropout_11 (Dropout)	(None, 4, 4, 32)	0
flatten_4 (Flatten)	(None, 512)	0
dense_7 (Dense)	(None, 512)	262656
dropout_12 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 10)	5130

8. Compile the Model

```
In [37]: # compile the model
model.compile(loss='categorical_crossentropy', optimizer='Adam',
metrics=['accuracy'])
```

9. Train the Model

```
In [38]: from keras.callbacks import ModelCheckpoint

# train the model
checkpointer = ModelCheckpoint(filepath='model_weights.best.hdf5', verbose=1,
save_best_only=True)
hist = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,
validation_data=(x_valid, y_valid), callbacks=[checkpointer],
verbose=2, shuffle=True)
```

```
Train on 45000 samples, validate on 5000 samples
Epoch 1/100
- 15s - loss: 1.7659 - accuracy: 0.3441 - val_loss: 1.4642 - val_accuracy: 0.4850
Epoch 0001: val_loss improved from 1.4642 to 1.4642, saving model to model_weights.best.hdf5
Epoch 2/100
- 15s - loss: 1.4993 - accuracy: 0.4504 - val_loss: 1.3215 - val_accuracy: 0.5300
Epoch 0002: val_loss improved from 1.3215 to 1.3214, saving model to model_weights.best.hdf5
Epoch 3/100
- 15s - loss: 1.4107 - accuracy: 0.4857 - val_loss: 1.2507 - val_accuracy: 0.5604
Epoch 0003: val_loss improved from 1.2507 to 1.2507, saving model to model_weights.best.hdf5
Epoch 4/100
- 15s - loss: 1.3623 - accuracy: 0.5058 - val_loss: 1.1976 - val_accuracy: 0.5858
Epoch 0004: val_loss improved from 1.1976 to 1.1975, saving model to model_weights.best.hdf5
Epoch 5/100
- 16s - loss: 1.3085 - accuracy: 0.5284 - val_loss: 1.1358 - val_accuracy: 0.6094
Epoch 0005: val_loss improved from 1.1358 to 1.1353, saving model to model_weights.best.hdf5
Epoch 6/100
- 15s - loss: 1.2796 - accuracy: 0.5359 - val_loss: 1.1268 - val_accuracy: 0.6048
Epoch 0006: val_loss improved from 1.1268 to 1.1267, saving model to model_weights.best.hdf5
Epoch 7/100
- 15s - loss: 1.2493 - accuracy: 0.5507 - val_loss: 1.1171 - val_accuracy: 0.6106
Epoch 0007: val_loss improved from 1.1171 to 1.1175, saving model to model_weights.best.hdf5
Epoch 8/100
- 15s - loss: 1.2303 - accuracy: 0.5562 - val_loss: 1.0861 - val_accuracy: 0.6306
Epoch 0008: val_loss improved from 1.0861 to 1.0860, saving model to model_weights.best.hdf5
Epoch 9/100
- 15s - loss: 1.2055 - accuracy: 0.5691 - val_loss: 1.0619 - val_accuracy: 0.6426
Epoch 0009: val_loss improved from 1.0619 to 1.0619, saving model to model_weights.best.hdf5
Epoch 10/100
- 15s - loss: 1.1855 - accuracy: 0.5734 - val_loss: 1.0209 - val_accuracy: 0.6498
Epoch 0010: val_loss improved from 1.0209 to 1.0208, saving model to model_weights.best.hdf5
Epoch 11/100
- 15s - loss: 1.1649 - accuracy: 0.5828 - val_loss: 1.0192 - val_accuracy: 0.6540
Epoch 0011: val_loss improved from 1.0192 to 1.0192, saving model to model_weights.best.hdf5
Epoch 12/100
- 15s - loss: 1.1558 - accuracy: 0.5875 - val_loss: 0.9984 - val_accuracy: 0.6636
Epoch 0012: val_loss improved from 0.9984 to 0.9983, saving model to model_weights.best.hdf5
Epoch 13/100
- 15s - loss: 1.1404 - accuracy: 0.5926 - val_loss: 0.9733 - val_accuracy: 0.6748
Epoch 0013: val_loss did not improve from 0.9733
Epoch 14/100
- 15s - loss: 1.1339 - accuracy: 0.5970 - val_loss: 0.9588 - val_accuracy: 0.6736
Epoch 0014: val_loss improved from 0.9588 to 0.9587, saving model to model_weights.best.hdf5
Epoch 15/100
- 15s - loss: 1.1161 - accuracy: 0.5995 - val_loss: 0.9537 - val_accuracy: 0.6636
Epoch 0015: val_loss improved from 0.9537 to 0.9537, saving model to model_weights.best.hdf5
Epoch 16/100
- 15s - loss: 1.1044 - accuracy: 0.6055 - val_loss: 0.9636 - val_accuracy: 0.6724
Epoch 0016: val_loss did not improve from 0.9537
Epoch 17/100
- 15s - loss: 1.0963 - accuracy: 0.6110 - val_loss: 0.9344 - val_accuracy: 0.6722
Epoch 0017: val_loss improved from 0.9344 to 0.9346, saving model to model_weights.best.hdf5
Epoch 18/100
- 15s - loss: 1.0893 - accuracy: 0.6108 - val_loss: 0.9407 - val_accuracy: 0.6758
Epoch 0018: val_loss did not improve from 0.9346
Epoch 19/100
- 15s - loss: 1.0830 - accuracy: 0.6126 - val_loss: 0.9042 - val_accuracy: 0.6972
Epoch 0019: val_loss improved from 0.9042 to 0.9041, saving model to model_weights.best.hdf5
Epoch 20/100
- 15s - loss: 1.0723 - accuracy: 0.6184 - val_loss: 0.9095 - val_accuracy: 0.6880
Epoch 0020: val_loss did not improve from 0.9041
Epoch 21/100
- 15s - loss: 1.0640 - accuracy: 0.6202 - val_loss: 0.9240 - val_accuracy: 0.6886
Epoch 0021: val_loss did not improve from 0.9041
Epoch 22/100
- 15s - loss: 1.0563 - accuracy: 0.6226 - val_loss: 0.9013 - val_accuracy: 0.6946
Epoch 0022: val_loss improved from 0.9016 to 0.9016, saving model to model_weights.best.hdf5
Epoch 23/100
- 15s - loss: 1.0523 - accuracy: 0.6220 - val_loss: 0.8850 - val_accuracy: 0.7118
Epoch 0023: val_loss improved from 0.9016 to 0.8849, saving model to model_weights.best.hdf5
Epoch 24/100
- 15s - loss: 1.0484 - accuracy: 0.6268 - val_loss: 0.8535 - val_accuracy: 0.7204
Epoch 0024: val_loss improved from 0.8849 to 0.8546, saving model to model_weights.best.hdf5
Epoch 25/100
- 16s - loss: 1.0414 - accuracy: 0.6289 - val_loss: 0.8651 - val_accuracy: 0.7080
Epoch 0025: val_loss did not improve from 0.8546
Epoch 26/100
- 15s - loss: 1.0329 - accuracy: 0.6333 - val_loss: 0.9203 - val_accuracy: 0.6950
Epoch 0026: val_loss did not improve from 0.8546
Epoch 27/100
- 15s - loss: 1.0339 - accuracy: 0.6316 - val_loss: 0.9300 - val_accuracy: 0.6844
Epoch 0027: val_loss did not improve from 0.8546
Epoch 28/100
- 15s - loss: 1.0235 - accuracy: 0.6342 - val_loss: 0.9007 - val_accuracy: 0.6886
Epoch 0028: val_loss did not improve from 0.8546
Epoch 29/100
- 15s - loss: 1.0260 - accuracy: 0.6352 - val_loss: 0.8611 - val_accuracy: 0.7102
Epoch 0029: val_loss did not improve from 0.8546
Epoch 30/100
- 15s - loss: 1.0194 - accuracy: 0.6383 - val_loss: 0.8520 - val_accuracy: 0.7074
Epoch 0030: val_loss improved from 0.8546 to 0.8519, saving model to model_weights.best.hdf5
Epoch 31/100
- 15s - loss: 1.0066 - accuracy: 0.6416 - val_loss: 0.8477 - val_accuracy: 0.7200
Epoch 0031: val_loss improved from 0.8519 to 0.8477, saving model to model_weights.best.hdf5
Epoch 32/100
- 15s - loss: 1.0061 - accuracy: 0.6430 - val_loss: 0.8421 - val_accuracy: 0.7138
Epoch 0032: val_loss improved from 0.8477 to 0.8420, saving model to model_weights.best.hdf5
Epoch 33/100
- 15s - loss: 1.0001 - accuracy: 0.6456 - val_loss: 0.8314 - val_accuracy: 0.7298
Epoch 0033: val_loss improved from 0.8420 to 0.8314, saving model to model_weights.best.hdf5
Epoch 34/100
- 15s - loss: 0.9941 - accuracy: 0.6470 - val_loss: 0.8352 - val_accuracy: 0.7150
Epoch 0034: val_loss did not improve from 0.8314
Epoch 35/100
- 15s - loss: 0.9940 - accuracy: 0.6471 - val_loss: 0.8256 - val_accuracy: 0.7218
Epoch 0035: val_loss improved from 0.8314 to 0.8256, saving model to model_weights.best.hdf5
Epoch 36/100
- 15s - loss: 0.9931 - accuracy: 0.6489 - val_loss: 0.8203 - val_accuracy: 0.7286
Epoch 0036: val_loss improved from 0.8256 to 0.8209, saving model to model_weights.best.hdf5
Epoch 37/100
- 15s - loss: 0.9868 - accuracy: 0.6503 - val_loss: 0.8679 - val_accuracy: 0.6992
Epoch 0037: val_loss did not improve from 0.8209
Epoch 38/100
- 15s - loss: 0.9885 - accuracy: 0.6489 - val_loss: 0.7887 - val_accuracy: 0.7384
Epoch 0038: val_loss improved from 0.8209 to 0.7886, saving model to model_weights.best.hdf5
Epoch 39/100
- 15s - loss: 0.9778 - accuracy: 0.6521 - val_loss: 0.8310 - val_accuracy: 0.7198
Epoch 0039: val_loss did not improve from 0.7886
Epoch 40/100
- 15s - loss: 0.9789 - accuracy: 0.6536 - val_loss: 0.7944 - val_accuracy: 0.7346
Epoch 0040: val_loss did not improve from 0.7886
Epoch 41/100
- 15s - loss: 0.9757 - accuracy: 0.6544 - val_loss: 0.8277 - val_accuracy: 0.7292
Epoch 0041: val_loss did not improve from 0.7886
Epoch 42/100
- 15s - loss: 0.9750 - accuracy: 0.6538 - val_loss: 0.8148 - val_accuracy: 0.7294
Epoch 0042: val_loss did not improve from 0.7886
Epoch 43/100
- 15s - loss: 0.9688 - accuracy: 0.6592 - val_loss: 0.7873 - val_accuracy: 0.7388
Epoch 0043: val_loss improved from 0.7886 to 0.7873, saving model to model_weights.best.hdf5
Epoch 44/100
- 15s - loss: 0.9636 - accuracy: 0.6585 - val_loss: 0.8219 - val_accuracy: 0.7274
Epoch 0044: val_loss did not improve from 0.7873
Epoch 45/100
- 15s - loss: 0.9569 - accuracy: 0.6607 - val_loss: 0.8037 - val_accuracy: 0.7348
Epoch 0045: val_loss did not improve from 0.7873
Epoch 46/100
- 16s - loss: 0.9562 - accuracy: 0.6612 - val_loss: 0.8097 - val_accuracy: 0.7332
Epoch 0046: val_loss did not improve from 0.7873
Epoch 47/100
- 15s - loss: 0.9571 - accuracy: 0.6602 - val_loss: 0.7919 - val_accuracy: 0.7386
Epoch 0047: val_loss did not improve from 0.7873
Epoch 48/100
- 15s - loss: 0.9570 - accuracy: 0.6583 - val_loss: 0.7752 - val_accuracy: 0.7442
Epoch 0048: val_loss improved from 0.7873 to 0.7751, saving model to model_weights.best.hdf5
Epoch 49/100
- 15s - loss: 0.9554 - accuracy: 0.6630 - val_loss: 0.7976 - val_accuracy: 0.7386
Epoch 0049: val_loss did not improve from 0.7751
Epoch 50/100
- 15s - loss: 0.9522 - accuracy: 0.6635 - val_loss: 0.8018 - val_accuracy: 0.7258
Epoch 0050: val_loss did not improve from 0.7751
Epoch 51/100
- 15s - loss: 0.9528 - accuracy: 0.6648 - val_loss: 0.7826 - val_accuracy: 0.7360
Epoch 0051: val_loss did not improve from 0.7751
Epoch 52/100
- 15s - loss: 0.9503 - accuracy: 0.6614 - val_loss: 0.8125 - val_accuracy: 0.7260
Epoch 0052: val_loss did not improve from 0.7751
Epoch 53/100
- 15s - loss: 0.9450 - accuracy: 0.6678 - val_loss: 0.7805 - val_accuracy: 0.7364
Epoch 0053: val_loss did not improve from 0.7751
Epoch 54/100
- 15s - loss: 0.9470 - accuracy: 0.6645 - val_loss: 0.7783 - val_accuracy: 0.7396
Epoch 0054: val_loss did not improve from 0.7751
Epoch 55/100
- 15s - loss: 0.9504 - accuracy: 0.6612 - val_loss: 0.8476 - val_accuracy: 0.7076
Epoch 0055: val_loss did not improve from 0.7751
Epoch 56/100
- 15s - loss: 0.9426 - accuracy: 0.6667 - val_loss: 0.8012 - val_accuracy: 0.7304
Epoch 0056: val_loss did not improve from 0.7751
Epoch 57/100
- 15s - loss: 0.9373 - accuracy: 0.6677 - val_loss: 0.7725 - val_accuracy: 0.7418
Epoch 0057: val_loss improved from 0.7751 to 0.7724, saving model to model_weights.best.hdf5
Epoch 58/100
- 15s - loss: 0.9489 - accuracy: 0.6649 - val_loss: 0.7760 - val_accuracy: 0.7492
Epoch 0058: val_loss did not improve from 0.7724
Epoch 59/100
- 15s - loss: 0.9386 - accuracy: 0.6679 - val_loss: 0.7761 - val_accuracy: 0.7412
Epoch 0059: val_loss did not improve from 0.7724
Epoch 60/100
- 15s - loss: 0.9370 - accuracy: 0.6704 - val_loss: 0.7905 - val_accuracy: 0.7316
Epoch 0060: val_loss did not improve from 0.7724
Epoch 61/100
- 15s - loss: 0.9339 - accuracy: 0.6708 - val_loss: 0.7938 - val_accuracy: 0.7382
Epoch 0061: val_loss did not improve from 0.7724
Epoch 62/100
- 15s - loss: 0.9325 - accuracy: 0.6685 - val_loss: 0.7621 - val_accuracy: 0.7374
Epoch 0062: val_loss improved from 0.7724 to 0.7621, saving model to model_weights.best.hdf5
Epoch 63/100
- 15s - loss: 0.9333 - accuracy: 0.6703 - val_loss: 0.7544 - val_accuracy: 0.7450
Epoch 0063: val_loss improved from 0.7621 to 0.7543, saving model to model_weights.best.hdf5
Epoch 64/100
- 15s - loss: 0.9293 - accuracy: 0.6724 - val_loss: 0.7948 - val_accuracy: 0.7320
Epoch 0064: val_loss did not improve from 0.7543
Epoch 65/100
- 15s - loss: 0.9320 - accuracy: 0.6704 - val_loss: 0.7691 - val_accuracy: 0.7384
Epoch 0065: val_loss did not improve from 0.7543
Epoch 66/100
- 15s - loss: 0.9304 - accuracy: 0.6705 - val_loss: 0.7641 - val_accuracy: 0.7416
Epoch 0066: val_loss did not improve from 0.7543
Epoch 67/100
- 15s - loss: 0.9255 - accuracy: 0.6724 - val_loss: 0.7989 - val_accuracy: 0.7300
Epoch 0067: val_loss did not improve from 0.7543
Epoch 68/100
- 15s - loss: 0.9234 - accuracy: 0.6725 - val_loss: 0.7773 - val_accuracy: 0.7314
Epoch 0068: val_loss did not improve from 0.7543
Epoch 69/100
- 15s - loss: 0.9130 - accuracy: 0.6760 - val_loss: 0.7832 - val_accuracy: 0.7328
Epoch 0069: val_loss did not improve from 0.7543
Epoch 70/100
- 15s - loss: 0.9191 - accuracy: 0.6776 - val_loss: 0.7564 - val_accuracy: 0.7416
Epoch 0070: val_loss did not improve from 0.7543
Epoch 71/100
- 15s - loss: 0.9220 - accuracy: 0.6743 - val_loss: 0.7744 - val_accuracy: 0.7342
Epoch 0071: val_loss did not improve from 0.7543
Epoch 72/100
- 15s - loss: 0.9202 - accuracy: 0.6761 - val_loss: 0.7836 - val_accuracy: 0.7378
Epoch 0072: val_loss did not improve from 0.7543
Epoch 73/100
- 15s - loss: 0.9149 - accuracy: 0.6777 - val_loss: 0.7749 - val_accuracy: 0.7348
Epoch 0073: val_loss did not improve from 0.7543
Epoch 74/100
- 15s - loss: 0.9208 - accuracy: 0.6737 - val_loss: 0.7489 - val_accuracy: 0.7404
Epoch 0074: val_loss improved from 0.7543 to 0.7489, saving model to model_weights.best.hdf5
Epoch 75/100
- 15s - loss: 0.9274 - accuracy: 0.6749 - val_loss: 0.8012 - val_accuracy: 0.7272
Epoch 0075: val_loss did not improve from 0.7489
Epoch 76/100
- 15s - loss: 0.9163 - accuracy: 0.6761 - val_loss: 0.8335 - val_accuracy: 0.7138
Epoch 0076: val_loss did not improve from 0.7489
Epoch 77/100
- 15s - loss: 0.9141 - accuracy: 0.6752 - val_loss: 0.7490 - val_accuracy: 0.7530
Epoch 0077: val_loss did not improve from 0.7489
Epoch 78/100
- 15s - loss: 0.9147 - accuracy: 0.6770 - val_loss: 0.7363 - val_accuracy: 0.7420
Epoch 0078: val_loss did not improve from 0.7489
Epoch 79/100
- 15s - loss: 0.9109 - accuracy: 0.6788 - val_loss: 0.7487 - val_accuracy: 0.7496
Epoch 0079: val_loss improved from 0.7489 to 0.7486, saving model to model_weights.best.hdf5
Epoch 80/100
- 15s - loss: 0.9076 - accuracy: 0.6806 - val_loss: 0.8131 - val_accuracy: 0.7250
Epoch 0080: val_loss did not improve from 0.7486
Epoch 81/100
- 15s - loss: 0.9087 - accuracy: 0.6787 - val_loss: 0.7461 - val_accuracy: 0.7486
Epoch 0081: val_loss improved from 0.7486 to 0.7461, saving model to model_weights.best.hdf5
Epoch 82/100
- 15s - loss: 0.9033 - accuracy: 0.6798 - val_loss: 0.7749 - val_accuracy: 0.7440
Epoch 0082: val_loss did not improve from 0.7461
Epoch 83/100
- 15s - loss: 0.9037 - accuracy: 0.6817 - val_loss: 0.7356 - val_accuracy: 0.7494
Epoch 0083: val_loss improved from 0.7461 to 0.7356, saving model to model_weights.best.hdf5
Epoch 84/100
- 15s - loss: 0.9070 - accuracy: 0.6794 - val_loss: 0.7747 - val_accuracy: 0.7416
Epoch 0084: val_loss did not improve from 0.7356
Epoch 85/100
- 15s - loss: 0.8951 - accuracy: 0.6787 - val_loss: 0.7298 - val_accuracy: 0.7598
Epoch 0085: val_loss improved from 0.7356 to 0.7298, saving model to model_weights.best.hdf5
Epoch 86/100
- 15s - loss: 0.9014 - accuracy: 0.6767 - val_loss: 0.7446 - val_accuracy: 0.7536
Epoch 0086: val_loss did not improve from 0.7298
Epoch 87/100
- 16s - loss: 0.9031 - accuracy: 0.6809 - val_loss: 0.7392 - val_accuracy: 0.7502
Epoch 0087: val_loss did not improve from 0.7298
Epoch 88/100
- 15s - loss: 0.9066 - accuracy: 0.6805 - val_loss: 0.7148 - val_accuracy: 0.7630
Epoch 0088: val_loss improved from 0.7298 to 0.7148, saving model to model_weights.best.hdf5
Epoch 89/100
- 15s - loss: 0.9099 - accuracy: 0.6787 - val_loss: 0.7368 - val_accuracy: 0.7514
Epoch 0089: val_loss did not improve from 0.7148
Epoch 90/100
- 15s - loss: 0.9005 - accuracy: 0.6806 - val_loss: 0.7359 - val_accuracy: 0.7540
Epoch 0090: val_loss did not improve from 0.7148
Epoch 91/100
- 15s - loss: 0.9006 - accuracy: 0.6794 - val_loss: 0.7393 - val_accuracy: 0.7526
Epoch 0091: val_loss did not improve from 0.7148
Epoch 92/100
- 15s - loss: 0.8980 - accuracy: 0.6805 - val_loss: 0.7513 - val_accuracy: 0.7500
Epoch 0092: val_loss did not improve from 0.7148
Epoch 93/100
- 15s - loss: 0.9088 - accuracy: 0.6786 - val_loss: 0.7682 - val_accuracy: 0.7398
Epoch 0093: val_loss did not improve from 0.7148
Epoch 94/100
- 14s 1315s - loss: 0.9012 - accuracy: 0.6793 - val_loss: 0.7493 - val_accuracy: 0.7440
Epoch 0094: val_loss did not improve from 0.7148
Epoch 95/100
- 15s - loss: 0.8963 - accuracy: 0.6844 - val_loss: 0.7441 - val_accuracy: 0.7480
Epoch 0095: val_loss did not improve from 0.7148
Epoch 96/100
- 15s - loss: 0.8952 - accuracy: 0.6824 - val_loss: 0.7383 - val_accuracy: 0.7382
Epoch 0096: val_loss did not improve from 0.7148
Epoch 97/100
- 15s - loss: 0.8954 - accuracy: 0.6854 - val_loss: 0.7
```


How to interpret loss and accuracy?

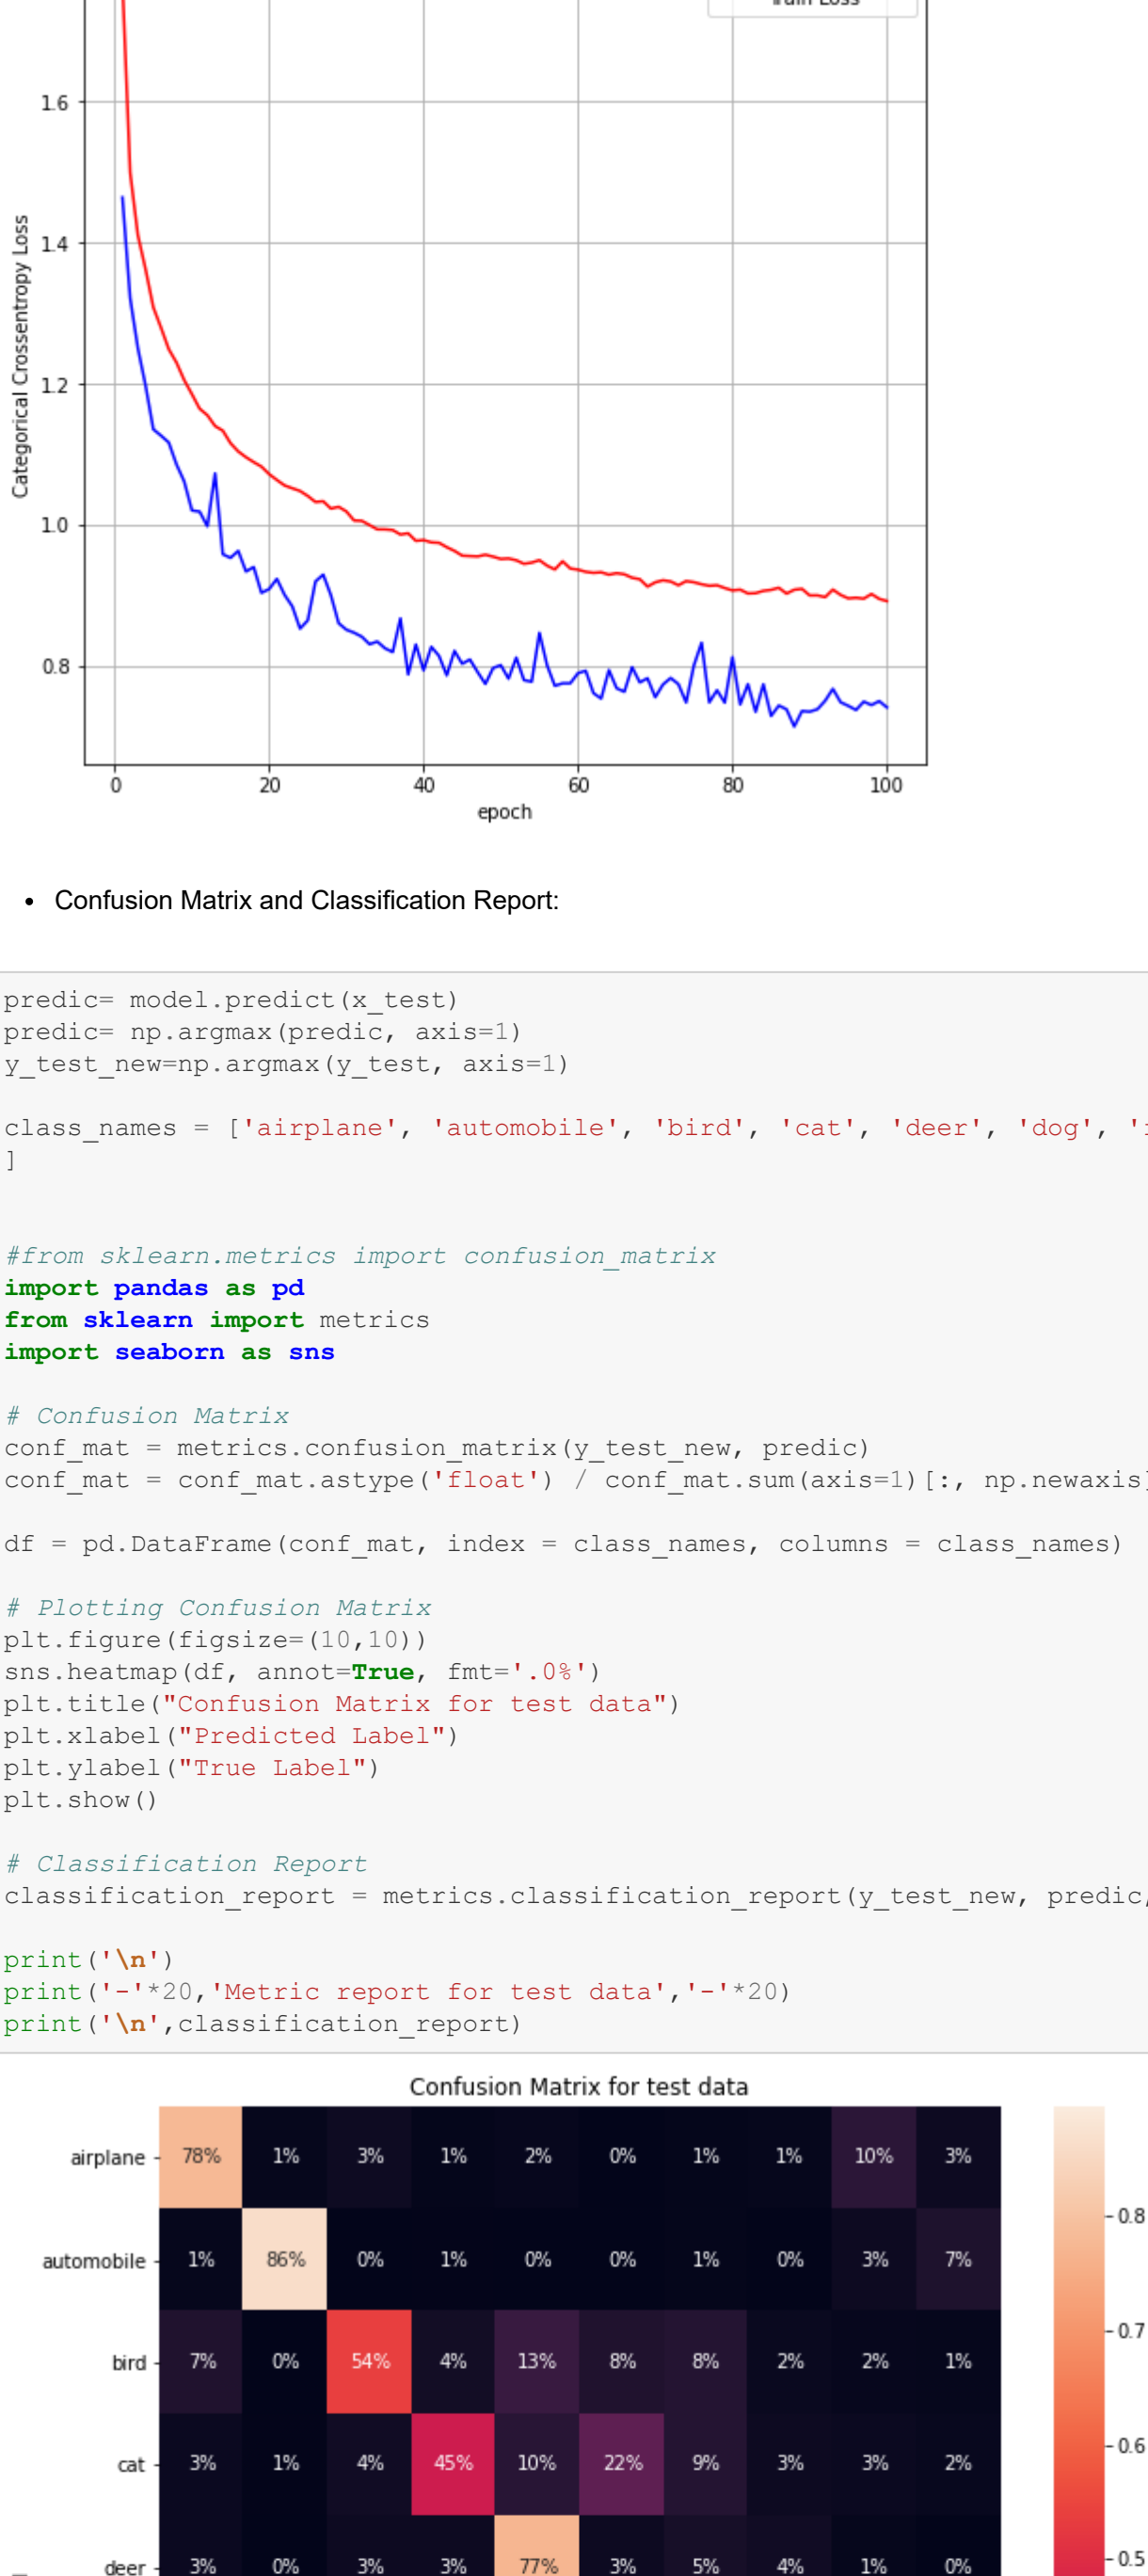
- Loss can be seen as a distance between the true values of the problem and the values predicted by the model.
- Greater the loss is, more huge is the errors made on the data.
- Accuracy can be seen as the number of errors made on the data.

That means :

- a low accuracy and huge loss means:
 - the model made huge errors on a lot of data
- a low accuracy but low loss means:
 - the model made little errors on a lot of data
- a great accuracy with low loss means:
 - the model made low errors on a few data (best case)

- The function below is used to update the plots for each epoch and error.

In [41]: <https://gist.github.com/greydanus/f6eee59ea1d90fcb3b534a25362cee4>



- Confusion Matrix and Classification Report:

In [42]:

```
predic=model.predict(x_test)
predic=np.argmax(predic,axis=1)
y_test_new=np.argmax(y_test,axis=1)

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

#from sklearn.metrics import confusion_matrix
import pandas as pd
from sklearn import metrics
import seaborn as sns

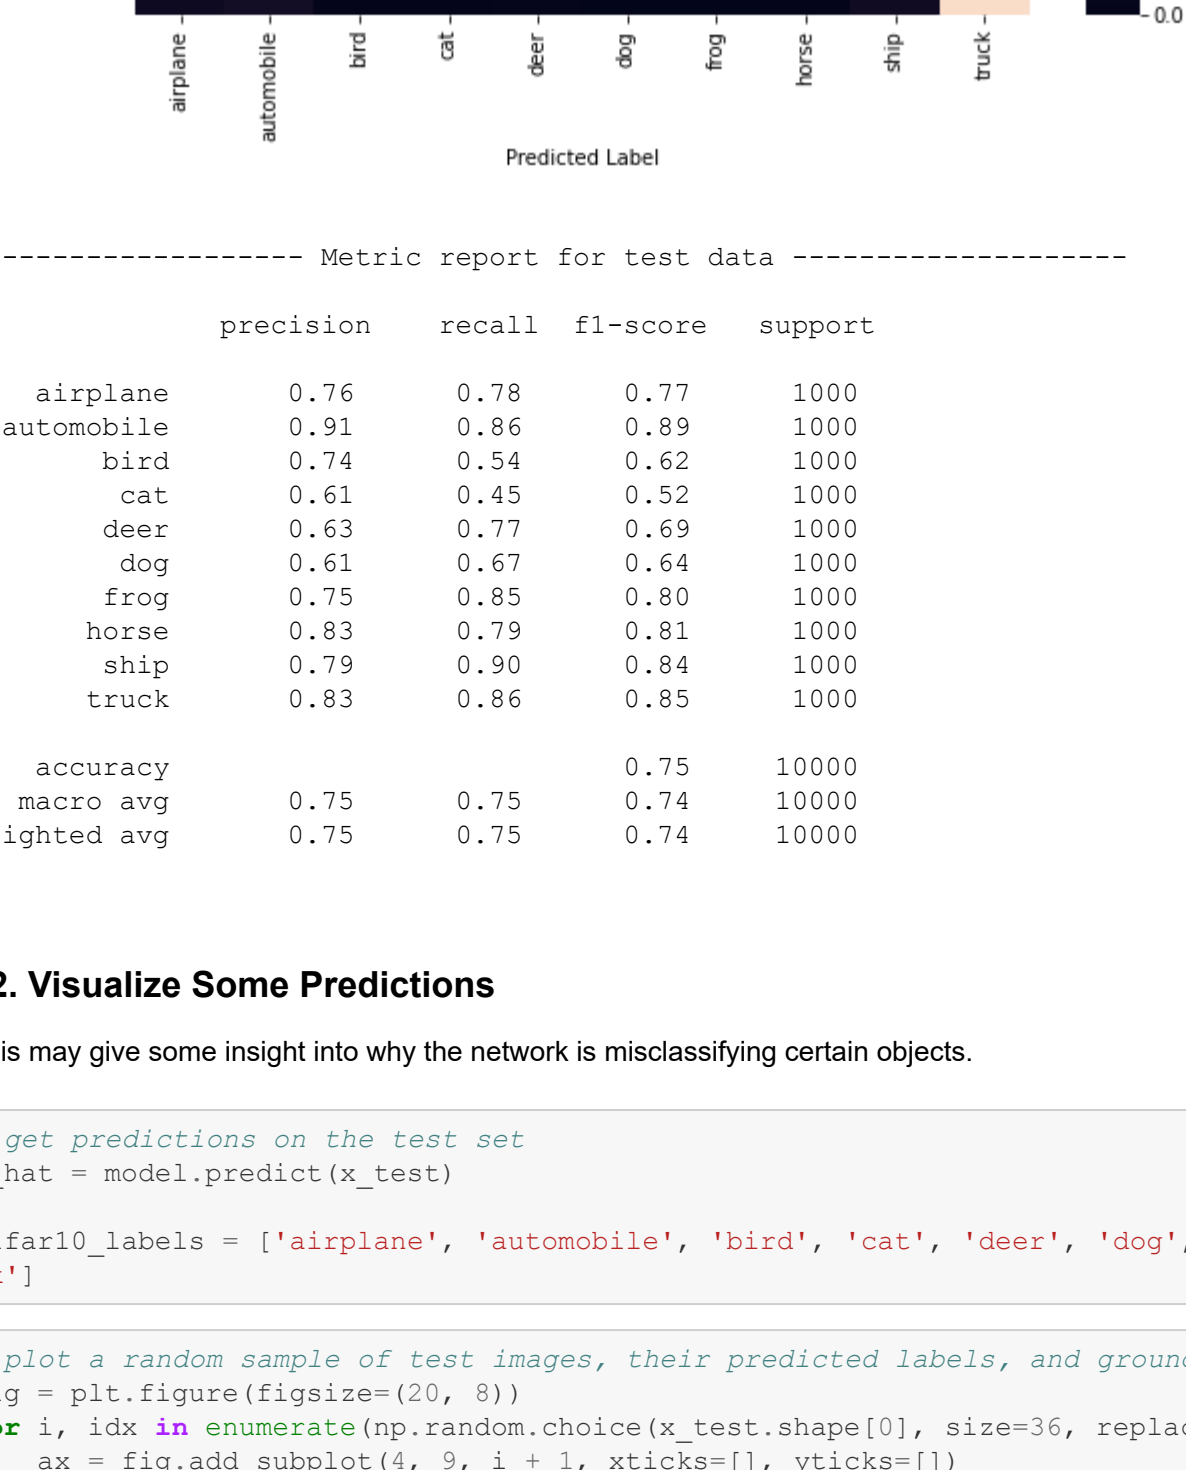
# Confusion Matrix
conf_mat = metrics.confusion_matrix(y_test_new, predic)
conf_mat = conf_mat.astype('float') / conf_mat.sum(axis=1)[:, np.newaxis]

df = pd.DataFrame(conf_mat, index = class_names, columns = class_names)

# Plotting Confusion Matrix
plt.figure(figsize=(10,10))
sns.heatmap(df, annot=True, fmt='.0%')
plt.title("Confusion Matrix for test data")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

# Classification Report
classification_report = metrics.classification_report(y_test_new, predic, target_names=class_names)

print('\n')
print('-'*20,'Metric report for test data','-'*20)
print('\n',classification_report)
```



12. Visualize Some Predictions

This may give some insight into why the network is misclassifying certain objects.

In [43]:

```
# get predictions on the test set
y_hat = model.predict(x_test)

cifar10_labels = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

In [44]:

```
# plot a random sample of test images, their predicted labels, and ground truth
fig = plt.figure(figsize=(20, 8))
for i, idx in enumerate(np.random.choice(x_test.shape[0], size=36, replace=False)):
    ax = fig.add_subplot(4, 9, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_test[idx]))
    pred_idx = np.argmax(y_hat[idx])
    true_idx = np.argmax(y_test[idx])
    ax.set_title("{} ({})*".format(cifar10_labels[pred_idx], cifar10_labels[true_idx]),
                  color="green" if pred_idx == true_idx else "red")
```



In [44]: