

# Convolutional Neural Networks

This notebook explains CNNs through the implementation in sklearn.

We train a CNN to classify images from the [CIFAR-10 database](#).

The CIFAR-10 dataset consists of 60,000 colour images of size (32x32) in 10 classes, with 6000 images per class.

- Here are the classes in the dataset, as well as 10 random images from each:
  - airplane
  - automobile
  - bird
  - cat
  - deer
  - dog
  - frog
  - horse
  - ship
  - truck

## 1. Load CIFAR-10 Database


```
In [95]: import keras
from keras.datasets import cifar10

# load the pre-shuffled train and test data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

## 2. Visualize Some Training Images

```
In [96]: import numpy as np
import matplotlib.pyplot as plt
matplotlib inline

fig = plt.figure(figsize=(20,5))
for i in range(40):
    ax = fig.add_subplot(4, 10, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_train[i]))
```



## 3. Rescale the Images by Dividing Every Pixel in Every Image by 255

- Let's see how a datapoint looks like:

```
In [97]: print(x_train[1])

[[[1154 177 187]
 [126 137 136]
 [105 104 95]
 ...
 [ 91 95 71]
 [ 87 90 71]
 [ 79 61 40]]

 [[140 160 169]
 [145 153 154]
 [125 125 118]
 ...
 [ 96 99 78]
 [ 77 80 62]
 [ 71 73 61]]

 [[140 155 164]
 [139 146 149]
 [115 115 112]
 ...
 [ 79 82 64]
 [ 68 70 55]
 [ 67 69 55]]

 ...

 [[175 167 166]
 [156 154 160]
 [154 160 170]
 ...
 [ 42 34 36]
 [ 61 53 57]
 [ 93 83 91]]

 [[165 154 128]
 [156 152 130]
 [159 161 142]
 ...
 [103 93 96]
 [123 114 120]
 [131 121 131]]

 [[163 148 120]
 [158 148 122]
 [163 156 133]
 ...
 [143 133 139]
 [143 134 142]
 [143 133 144]]]
```

- Let's normalize each cell value [0.255], so that each cell value lies within the same distribution [0,1].

- That is values change from [0.255] -> [0,1].

- We normalize using the formula:

$$\frac{(X - X_{\min})}{(X_{\max} - X_{\min})}$$

- Here,  $X_{\min}=0$  and  $X_{\max}=255$ .

- Hence,

$$\frac{(X - X_{\min})}{(X_{\max} - X_{\min})} = \frac{X}{255}$$

```
In [98]: x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255
```

## 4. Break Dataset into Training, Testing, and Validation Sets

- A datapoint in the dataset could belong to any one of the 10 classes mentioned earlier.
- Let's convert the output into one-hot encoded vector of size 10.
- This will be helpful for our CNN to recognize a given image datapoint into one of the 10 classes.

```
In [99]: from keras.utils import np_utils

# one-hot encoding
num_classes = len(np.unique(y_train))
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [100]: # break training set into training and validation sets
(x_train, x_valid) = x_train[:5000], x_train[5000:]
(y_train, y_valid) = y_train[:5000], y_train[5000:]

# print shape of training set
print('x_train shape:', x_train.shape)

# print number of training, validation, and test images
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
print(x_valid.shape[0], 'validation samples')
```

x\_train shape: (45000, 32, 32, 3)  
45000 train samples  
10000 test samples  
5000 validation samples

## 5. Data Augmentation

- Data augmentation helps in creating invariance in the data, especially in image data.
  - Invariance in data could be for:
    - Shift
    - Rotation
    - Zoom
    - Noise, etc

```
In [101]: from keras.preprocessing.image import ImageDataGenerator

# create and configure augmented image generator
datagen_train = ImageDataGenerator(
    width_shift_range=0.1, # randomly shift images horizontally (10% of total width)
    height_shift_range=0.1, # randomly shift images vertically (10% of total height)
    horizontal_flip=True) # randomly flip images horizontally

# fit augmented image generator on data
datagen_train.fit(x_train)
```

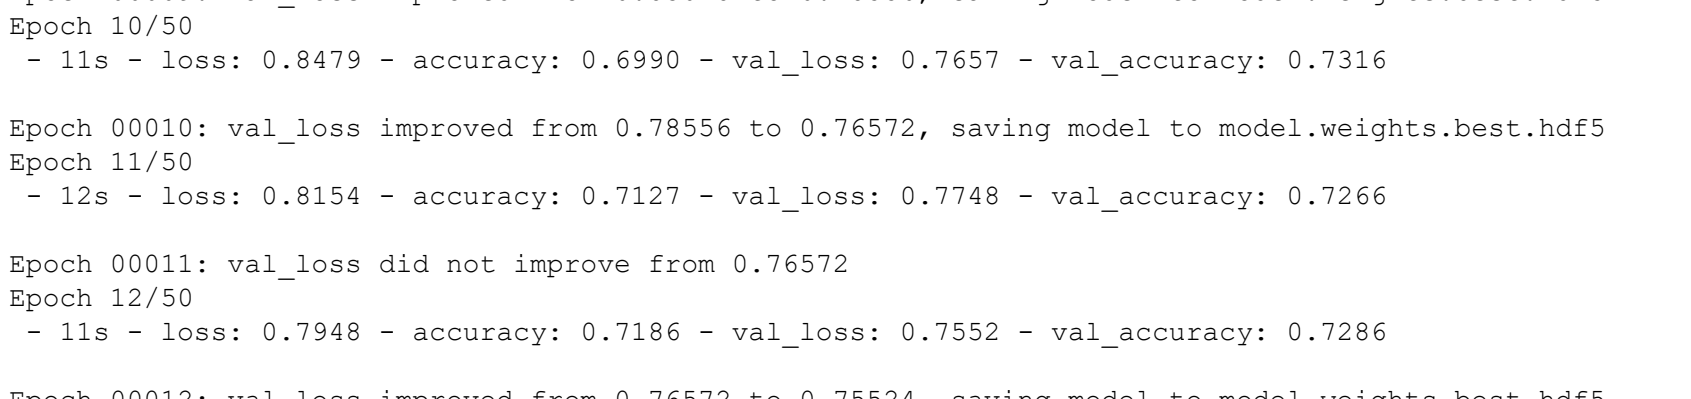
## 6. Visualize Original and Augmented Images

```
In [102]: import matplotlib.pyplot as plt

# take subset of training data
x_train_subset = x_train[1:12]

# visualize subset of training data
fig = plt.figure(figsize=(20,2))
for i in range(0, len(x_train_subset)):
    ax = fig.add_subplot(1, 12, i+1)
    ax.imshow(x_train_subset[i])
fig.suptitle('Subset of Original Training Images', fontsize=20)
plt.show()

# visualize augmented images
fig = plt.figure(figsize=(20,2))
for i in range(0, len(x_train_subset)):
    ax = fig.add_subplot(1, 12, i+1)
    ax.imshow(x_train_subset[i])
fig.suptitle('Augmented Images', fontsize=20)
plt.show()
break;
```



## 7. Define the Model Architecture

```
In [103]: from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

input_shape = x_train[1].shape

model = Sequential()
model.add(Conv2D(filters=32, kernel_size=3, padding='same', kernel_initializer='glorot_uniform',
    activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2, padding='same', kernel_initializer='glorot_uniform',
    activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.25))
model.add(Conv2D(filters=64, kernel_size=2, kernel_initializer='glorot_uniform', padding='same',
    activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(512, kernel_initializer='glorot_uniform', activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(10, activation='softmax'))

model.summary()
```

```
Epoch 00026: val_loss improved from 0.64494 to 0.64125, saving model
Epoch 27/50
 11s - loss: 0.5596 - accuracy: 0.8012 - val_loss: 0.6457 - val_acc: 0.6370
Epoch 00027: val_loss did not improve from 0.64125
Epoch 28/50
 12s - loss: 0.5444 - accuracy: 0.8072 - val_loss: 0.6456 - val_acc: 0.6370
Epoch 00028: val_loss did not improve from 0.64125
Epoch 29/50
 12s - loss: 0.5314 - accuracy: 0.8108 - val_loss: 0.6469 - val_acc: 0.6370
Epoch 00029: val_loss did not improve from 0.64125
Epoch 30/50
 12s - loss: 0.5244 - accuracy: 0.8146 - val_loss: 0.6829 - val_acc: 0.6370
Epoch 00030: val_loss did not improve from 0.64125
Epoch 31/50
 12s - loss: 0.5217 - accuracy: 0.8177 - val_loss: 0.6371 - val_acc: 0.6370
Epoch 00031: val_loss improved from 0.64125 to 0.63709, saving model
Epoch 32/50
 12s - loss: 0.5054 - accuracy: 0.8227 - val_loss: 0.6564 - val_acc: 0.6370
Epoch 00032: val_loss did not improve from 0.63709
Epoch 33/50
 11s - loss: 0.4973 - accuracy: 0.8234 - val_loss: 0.6528 - val_acc: 0.6370
Epoch 00033: val_loss did not improve from 0.63709
Epoch 34/50
 12s - loss: 0.4937 - accuracy: 0.8256 - val_loss: 0.6732 - val_acc: 0.6370
Epoch 00034: val_loss did not improve from 0.63709
```

## 8. Compile the Model

```
In [104]: # compile the model
model.compile(loss='categorical_crossentropy', optimizer='Adamax',
    metrics=['accuracy'])
```

## 9. Train the Model

**Note:**

- verbose=0 will show you nothing (silent)
- verbose=1 will show you an animated progress bar like this:
  - progres\_bar
- verbose=2 will just mention the number of epoch like this:
  - Epoch = 1/10

```
In [105]: from keras.callbacks import ModelCheckpoint

batch_size = 35
epochs = 50

# train the model
checkpointer = ModelCheckpoint(filepath='model.weights.best.hdf5', verbose=1,
    save_best_only=True)
hist = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,
    validation_data=(x_valid, y_valid), callbacks=[checkpointer],
    verbose=2, shuffle=True)
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/50  
- 12s - loss: 1.6698 - accuracy: 0.3874 - val\_loss: 1.3551 - val\_accuracy: 0.5184

Epoch 0001: val\_loss improved from inf to 1.35512, saving model to model.weights.best.hdf5

Epoch 2/50  
- 11s - loss: 1.3958 - accuracy: 0.5094 - val\_loss: 1.2253 - val\_accuracy: 0.5532

Epoch 0002: val\_loss improved from 1.35512 to 1.22526, saving model to model.weights.best.hdf5

Epoch 3/50  
- 11s - loss: 1.2383 - accuracy: 0.5564 - val\_loss: 1.1120 - val\_accuracy: 0.6148

Epoch 0003: val\_loss improved from 1.22526 to 1.11199, saving model to model.weights.best.hdf5

Epoch 4/50  
- 13s - loss: 1.1399 - accuracy: 0.5939 - val\_loss: 1.0113 - val\_accuracy: 0.6498

Epoch 0004: val\_loss improved from 1.11199 to 1.01130, saving model to model.weights.best.hdf5

Epoch 5/50  
- 12s - loss: 1.0773 - accuracy: 0.6150 - val\_loss: 0.9489 - val\_accuracy: 0.6730

Epoch 0005: val\_loss improved from 1.01130 to 0.94893, saving model to model.weights.best.hdf5

Epoch 6/50  
- 11s - loss: 1.0180 - accuracy: 0.6368 - val\_loss: 0.9235 - val\_accuracy: 0.6824

Epoch 0006: val\_loss improved from 0.94893 to 0.92349, saving model to model.weights.best.hdf5

Epoch 7/50  
- 12s - loss: 0.9618 - accuracy: 0.6578 - val\_loss: 0.8814 - val\_accuracy: 0.6976

Epoch 0007: val\_loss improved from 0.92349 to 0.88143, saving model to model.weights.best.hdf5

Epoch 8/50  
- 12s - loss: 0.9188 - accuracy: 0.6748 - val\_loss: 0.8303 - val\_accuracy: 0.7166

Epoch 0008: val\_loss improved from 0.88143 to 0.83028, saving model to model.weights.best.hdf5

Epoch 9/50  
- 12s - loss: 0.8825 - accuracy: 0.6869 - val\_loss: 0.7856 - val\_accuracy: 0.7230

Epoch 0009: val\_loss improved from 0.83028 to 0.78556, saving model to model.weights.best.hdf5

Epoch 10/50  
- 11s - loss: 0.8479 - accuracy: 0.6990 - val\_loss: 0.7657 - val\_accuracy: 0.7316

Epoch 0010: val\_loss improved from 0.78556 to 0.76572, saving model to model.weights.best.hdf5

Epoch 11/50  
- 12s - loss: 0.8154 - accuracy: 0.7127 - val\_loss: 0.7748 - val\_accuracy: 0.7266

Epoch 0011: val\_loss did not improve from 0.76572

Epoch 12/50  
- 11s - loss: 0.7948 - accuracy: 0.7186 - val\_loss: 0.7552 - val\_accuracy: 0.7286

Epoch 0012: val\_loss improved from 0.76572 to 0.75524, saving model to model.weights.best.hdf5

Epoch 13/50  
- 12s - loss: 0.7692 - accuracy: 0.7300 - val\_loss: 0.7202 - val\_accuracy: 0.7470

Epoch 0013: val\_loss improved from 0.75524 to 0.72023, saving model to model.weights.best.hdf5

Epoch 14/50  
- 12s - loss: 0.7475 - accuracy: 0.7370 - val\_loss: 0.7012 - val\_accuracy: 0.7582

Epoch 0014: val\_loss improved from 0.72023 to 0.70122, saving model to model.weights.best.hdf5

Epoch 15/50  
- 11s - loss: 0.7208 - accuracy: 0.7473 - val\_loss: 0.6869 - val\_accuracy: 0.7532

Epoch 0015: val\_loss improved from 0.70122 to 0.68690, saving model to model.weights.best.hdf5

Epoch 16/50  
- 11s - loss: 0.7075 - accuracy: 0.7519 - val\_loss: 0.6943 - val\_accuracy: 0.7598

Epoch 0016: val\_loss did not improve from 0.68690

Epoch 17/50  
- 11s - loss: 0.6866 - accuracy: 0.7577 - val\_loss: 0.6832 - val\_accuracy: 0.7566

Epoch 0017: val\_loss improved from 0.68690 to 0.68323, saving model to model.weights.best.hdf5

Epoch 18/50  
- 12s - loss: 0.6717 - accuracy: 0.7641 - val\_loss: 0.6934 - val\_accuracy: 0.7564

Epoch 0018: val\_loss did not improve from 0.68323

Epoch 19/50  
- 12s - loss: 0.6609 - accuracy: 0.7646 - val\_loss: 0.6757 - val\_accuracy: 0.7592

Epoch 0019: val\_loss improved from 0.68323 to 0.67565, saving model to model.weights.best.hdf5

Epoch 20/50  
- 12s - loss: 0.6402 - accuracy: 0.7747 - val\_loss: 0.6568 - val\_accuracy: 0.7698

Epoch 0020: val\_loss improved from 0.67565 to 0.65682, saving model to model.weights.best.hdf5

Epoch 21/50  
- 11s - loss: 0.6241 - accuracy: 0.7820 - val\_loss: 0.6528 - val\_accuracy: 0.7724

Epoch 0021: val\_loss improved from 0.65682 to 0.65277, saving model to model.weights.best.hdf5

Epoch 22/50  
- 12s - loss: 0.6135 - accuracy: 0.7819 - val\_loss: 0.6587 - val\_accuracy: 0.7722

Epoch 0022: val\_loss did not improve from 0.65277

Epoch 23/50  
- 11s - loss: 0.5938 - accuracy: 0.7906 - val\_loss: 0.6476 - val\_accuracy: 0.7708

Epoch 0023: val\_loss improved from 0.65277 to 0.64755, saving model to model.weights.best.hdf5

Epoch 24/50  
- 11s - loss: 0.5840 - accuracy: 0.7909 - val\_loss: 0.6595 - val\_accuracy: 0.7716

Epoch 0024: val\_loss did not improve from 0.64755

Epoch 25/50  
- 12s - loss: 0.5783 - accuracy: 0.7960 - val\_loss: 0.6449 - val\_accuracy: 0.7782

Epoch 0025: val\_loss improved from 0.64755 to 0.64494, saving model to model.weights.best.hdf5

Epoch 26/50  
- 12s - loss: 0.5642 - accuracy: 0.8024 - val\_loss: 0.6412 - val\_accuracy: 0.7820

Epoch 0026: val\_loss improved from 0.64494 to 0.64125, saving model to model.weights.best.hdf5

Epoch 27/50  
- 12s - loss: 0.5596 - accuracy: 0.8012 - val\_loss: 0.6457 - val\_accuracy: 0.7762

Epoch 0027: val\_loss did not improve from 0.64125

Epoch 28/50  
- 12s - loss: 0.5444 - accuracy: 0.8072 - val\_loss: 0.6456 - val\_accuracy: 0.7802

Epoch 0028: val\_loss did not improve from 0.64125

Epoch 29/50  
- 12s - loss: 0.5314 - accuracy: 0.8108 - val\_loss: 0.6469 - val\_accuracy: 0.7768

Epoch 0029: val\_loss did not improve from 0.64125

Epoch 30/50  
- 12s - loss: 0.5244 - accuracy: 0.8146 - val\_loss: 0.6829 - val\_accuracy: 0.7660

Epoch 0030: val\_loss did not improve from 0.64125

Epoch 31/50  
- 12s - loss: 0.5217 - accuracy: 0.8177 - val\_loss: 0.6371 - val\_accuracy: 0.7746

Epoch 0031: val\_loss improved from 0.64125 to 0.63709, saving model to model.weights.best.hdf5

Epoch 32/50  
- 12s - loss: 0.5054 - accuracy: 0.8227 - val\_loss: 0.6564 - val\_accuracy: 0.7754

Epoch 0032: val\_loss did not improve from 0.63709

Epoch 33/50  
- 11s - loss: 0.4973 - accuracy: 0.8234 - val\_loss: 0.6528 - val\_accuracy: 0.7752

Epoch 0033: val\_loss did not improve from 0.63709

Epoch 34/50  
- 12s - loss: 0.4937 - accuracy: 0.8256 - val\_loss: 0.6732 - val\_accuracy: 0.7614

Epoch 0034: val\_loss did not improve from 0.63709

Epoch 35/50  
- 11s - loss: 0.4790 - accuracy: 0.8314 - val\_loss: 0.6707 - val\_accuracy: 0.7760

Epoch 0035: val\_loss did not improve from 0.63709

Epoch 36/50  
- 12s - loss: 0.4761 - accuracy: 0.8297 - val\_loss: 0.6683 - val\_accuracy: 0.7784

Epoch 0036: val\_loss did not improve from 0.63709

Epoch 37/50  
- 12s - loss: 0.4722 - accuracy: 0.8314 - val\_loss: 0.6430 - val\_accuracy: 0.7792

Epoch 0037: val\_loss did not improve from 0.63709

Epoch 38/50  
- 12s - loss: 0.4566 - accuracy: 0.8389 - val\_loss: 0.6701 - val\_accuracy: 0.7778

Epoch 0038: val\_loss did not improve from 0.63709

Epoch 39/50  
- 12s - loss: 0.4588 - accuracy: 0.8354 - val\_loss: 0.6546 - val\_accuracy: 0.7772

Epoch 0039: val\_loss did not improve from 0.63709

Epoch 40/50  
- 12s - loss: 0.4496 - accuracy: 0.8404 - val\_loss: 0.6607 - val\_accuracy: 0.7786

Epoch 0040: val\_loss did not improve from 0.63709

Epoch 41/50  
- 11s - loss: 0.4422 - accuracy: 0.8429 - val\_loss: 0.6673 - val\_accuracy: 0.7742

Epoch 0041: val\_loss did not improve from 0.63709

Epoch 42/50  
- 12s - loss: 0.4399 - accuracy: 0.8430 - val\_loss: 0.6429 - val\_accuracy: 0.7818

Epoch 0042: val\_loss did not improve from 0.63709

Epoch 43/50  
- 11s - loss: 0.4357 - accuracy: 0.8452 - val\_loss: 0.6674 - val\_accuracy: 0.7714

Epoch 0043: val\_loss did not improve from 0.63709

Epoch 44/50  
- 12s - loss: 0.4275 - accuracy: 0.8467 - val\_loss: 0.6412 - val\_accuracy: 0.7850

Epoch 0044: val\_loss did not improve from 0.63709

Epoch 45/50  
- 11s - loss: 0.4216 - accuracy: 0.8509 - val\_loss: 0.6421 - val\_accuracy: 0.7864

Epoch 0045: val\_loss did not improve from 0.63709

Epoch 46/50  
- 12s - loss: 0.4149 - accuracy: 0.8515 - val\_loss: 0.6524 - val\_accuracy: 0.7812

Epoch 0046: val\_loss did not improve from 0.63709

Epoch 47/50  
- 12s - loss: 0.4085 - accuracy: 0.8542 - val\_loss: 0.6496 - val\_accuracy: 0.7834

Epoch 0047: val\_loss did not improve from 0.63709

Epoch 48/50  
- 11s - loss: 0.4032 - accuracy: 0.8569 - val\_loss: 0.6642 - val\_accuracy: 0.7832

Epoch 0048: val\_loss did not improve from 0.63709

Epoch 49/50  
- 11s - loss: 0.4013 - accuracy: 0.8577 - val\_loss: 0.6512 - val\_accuracy: 0.7872

Epoch 0049: val\_loss did not improve from 0.63709

Epoch 50/50  
- 11s - loss: 0.3975 - accuracy: 0.8580 - val\_loss: 0.6483 - val\_accuracy: 0.7844

Epoch 0050: val\_loss did not improve from 0.63709

## 10. Load the Model with the Best Validation Accuracy

```
In [106]: # load the weights that yielded the best validation accuracy
model.load_weights('model.weights.best.hdf5')
```

## 11. Calculate Classification Accuracy on Test Set

```
In [107]: # evaluate and print test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Test loss: 0.6550379812240601  
Test accuracy: 0.7741000056266785

## How to interpret loss and accuracy?

- Loss can be seen as a distance between the true values of the problem and the values predicted by the model.
  - Greater the loss is, more huge is the errors made on the data.
- Accuracy can be seen as the number of errors made on the data.

- That means:**
- a low accuracy and huge loss means:
    - the model made huge errors on a lot of data
  - a low accuracy but low loss means:
    - the model made little errors on a lot of data
  - a great accuracy with low loss means:
    - the model made low errors on a few data (best case)